



Navigating the shallow waters of Kubernetes Security

*Me harty's guide
fer all ye scallywags*



The modern application stack



Application



Orchestrator



Container

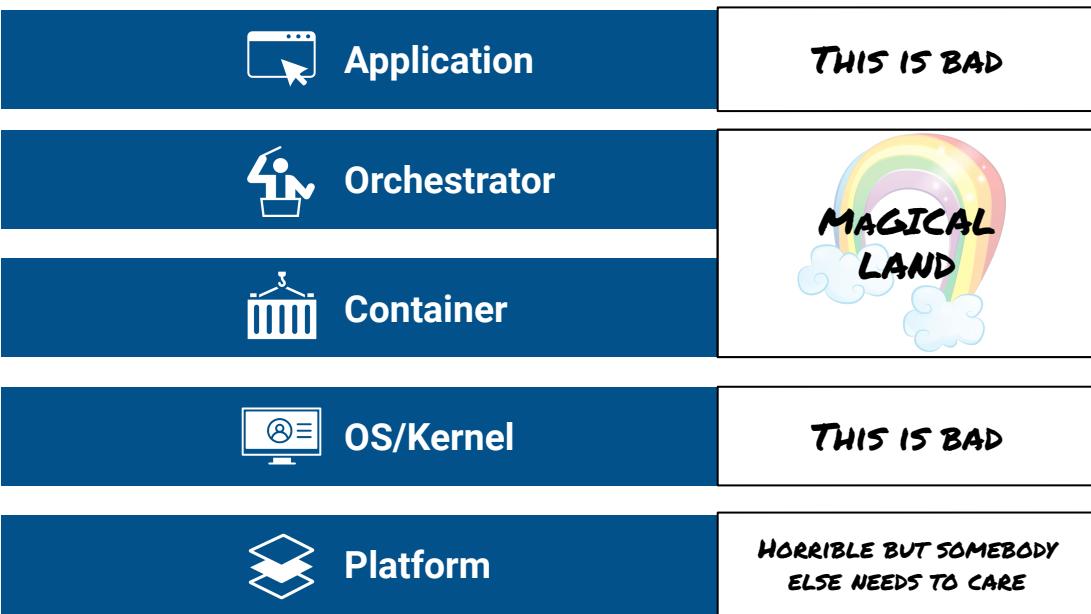


OS/Kernel



Platform

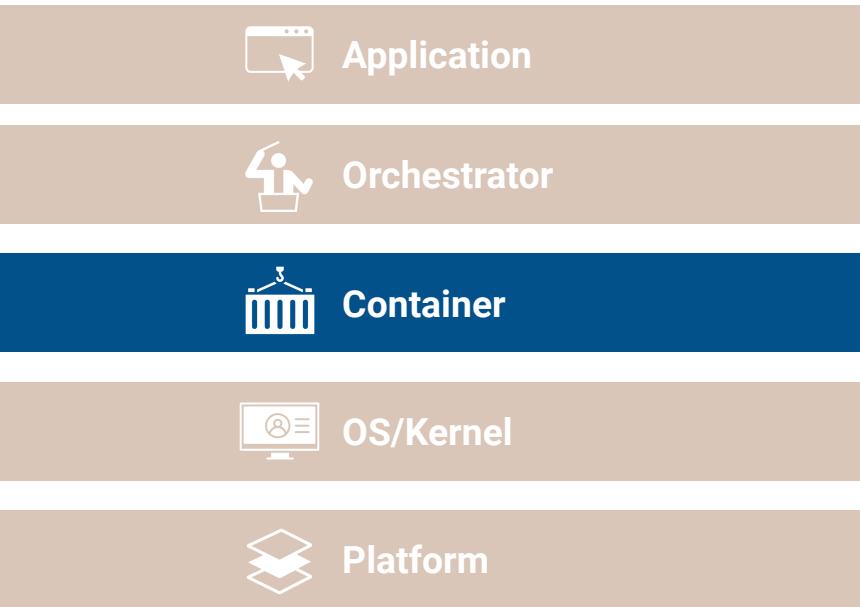
The modern application stack



Know yer watters



The modern application stack



Can containers help with my security?

Yes, if you build them, deploy and run them properly.

Types of issue a container can help with limiting the damage

- Persistence
- Tooling/living-of-the-land
- Path traversal
- Uncontrolled resource consumption/DoS



Where containers will not help

If your app is bad, it won't fix it

- Injection/Insecure Deserialization/RCE
- Use-after-free, TOCTOU, overflow
- XSS, CSRF, SSRF
- Change CRI defaults, get yourself in trouble



Get'er going



The modern application stack



Application



Orchestrator



Container



OS/Kernel



Platform

Infrastructure as a Code



Chef

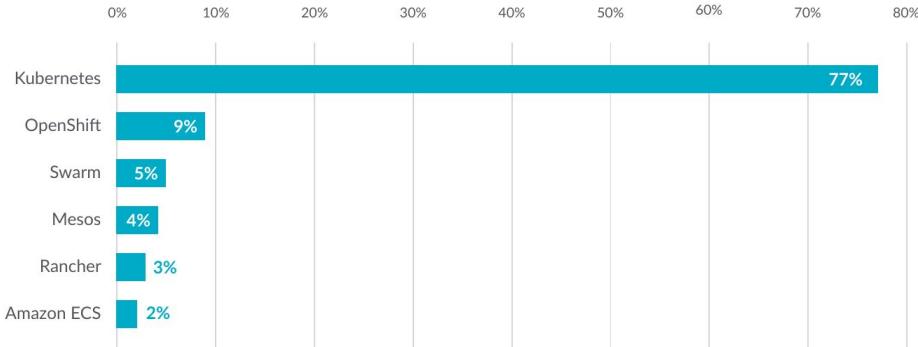


Puppet



Terraform

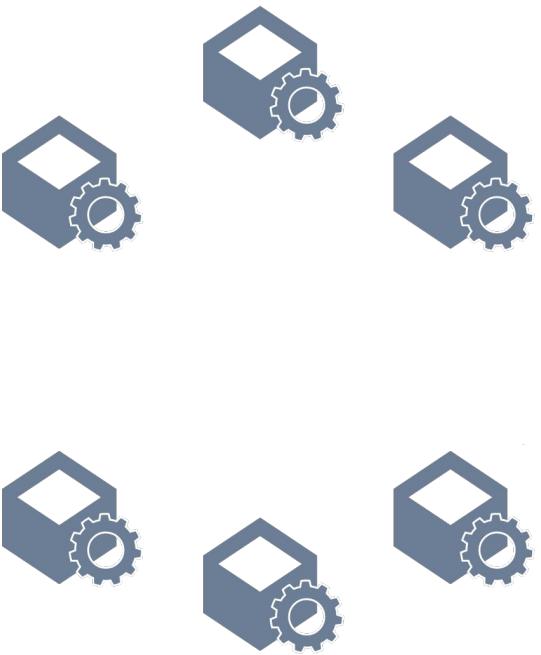
Kubernetes!

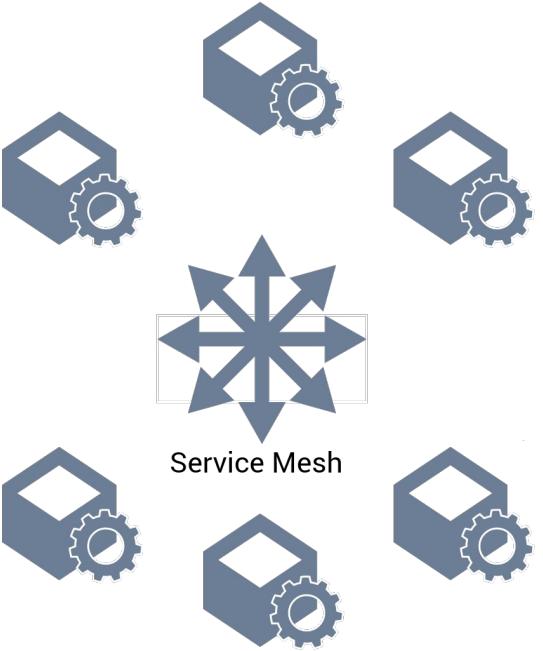


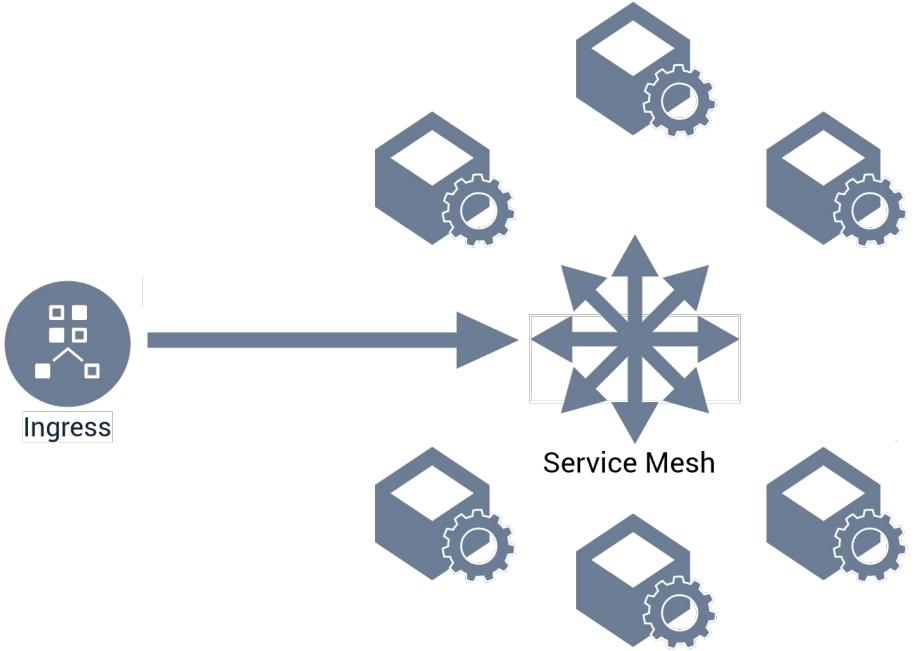
Core K8s concepts

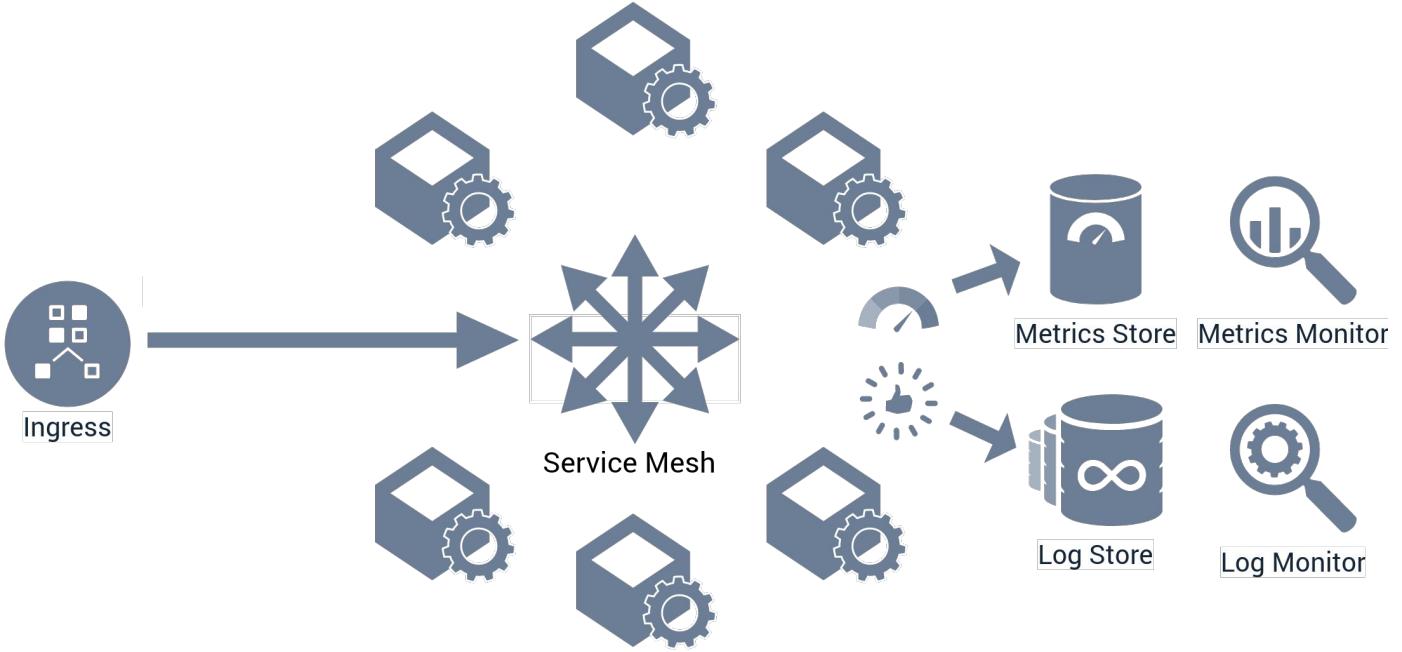
1. Declarative deployment
2. No hidden APIs
3. Meet the developer where they are at

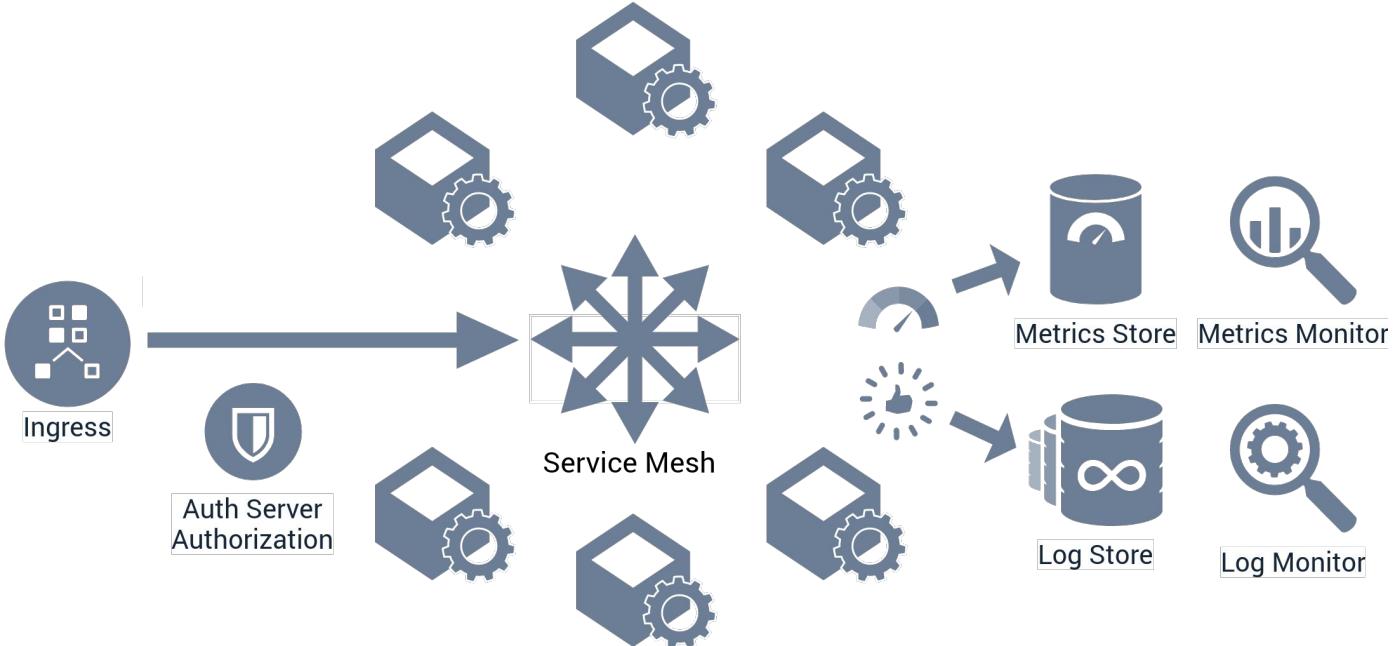


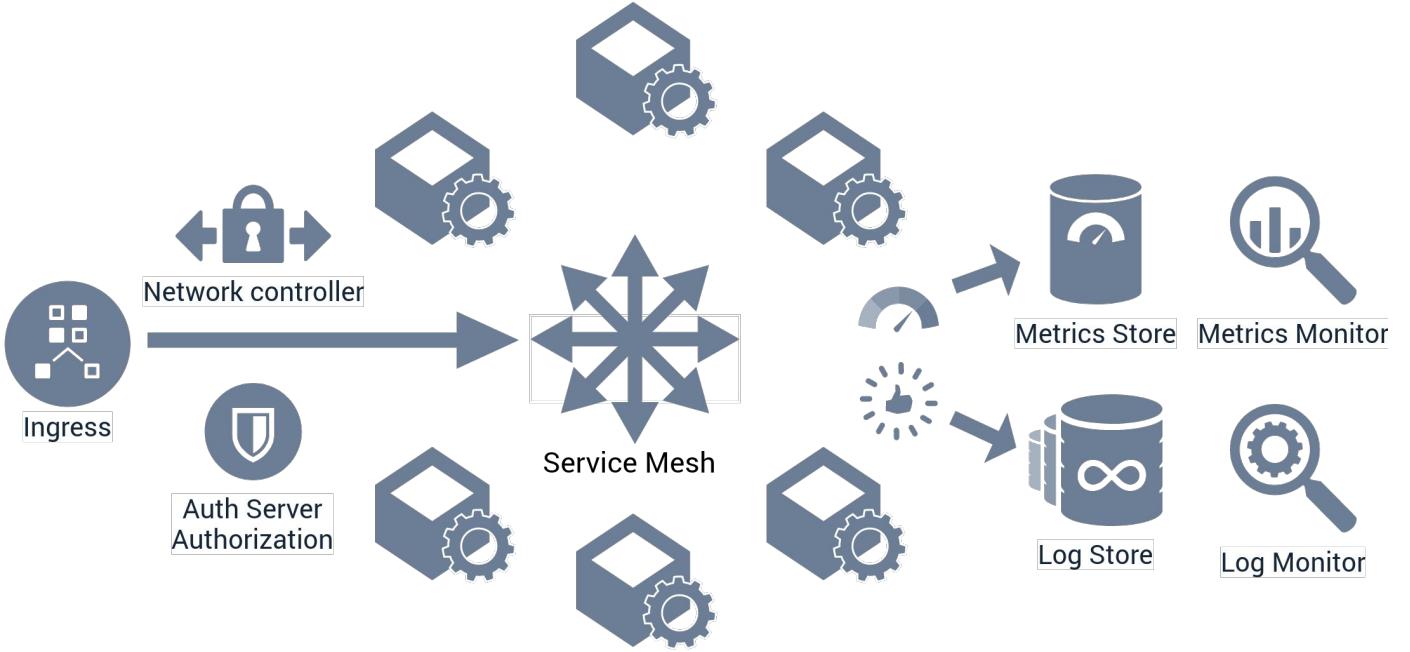


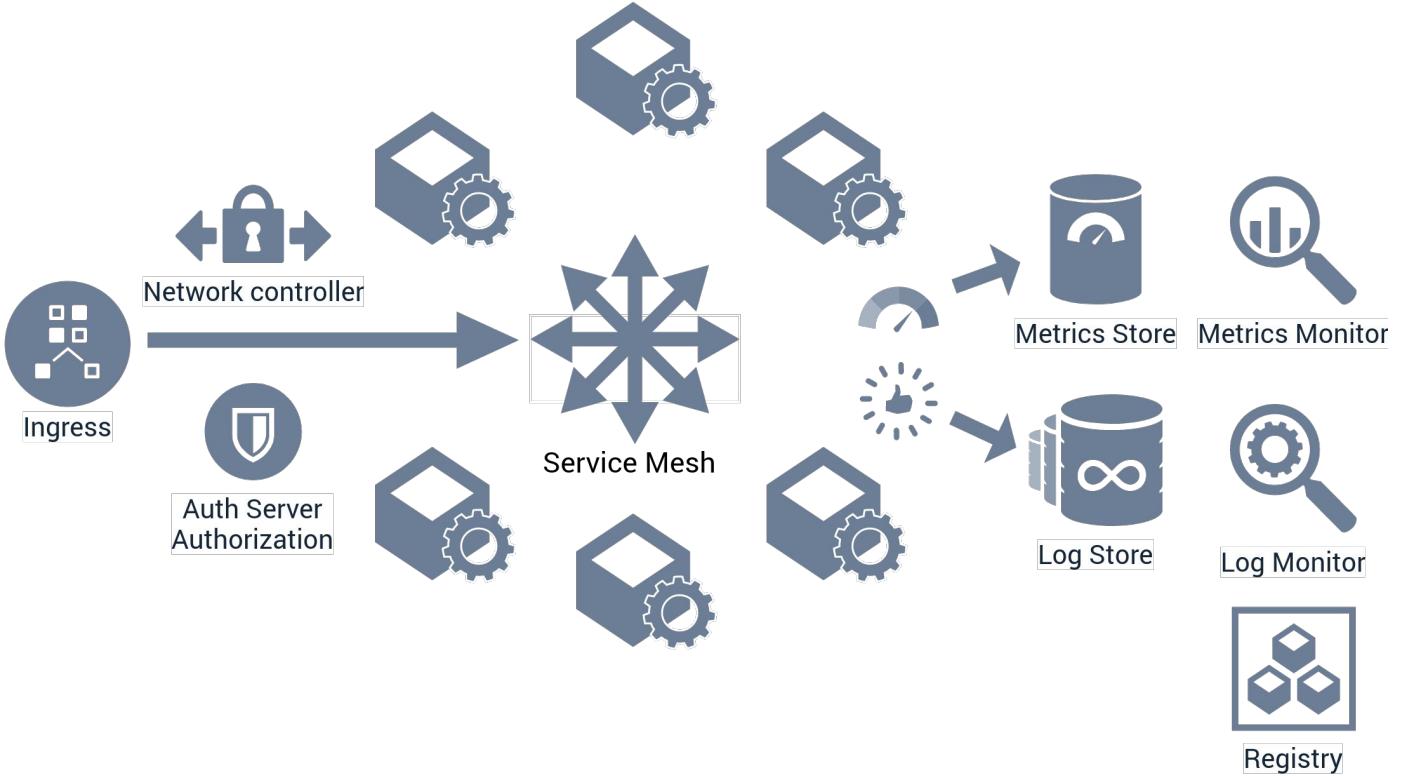


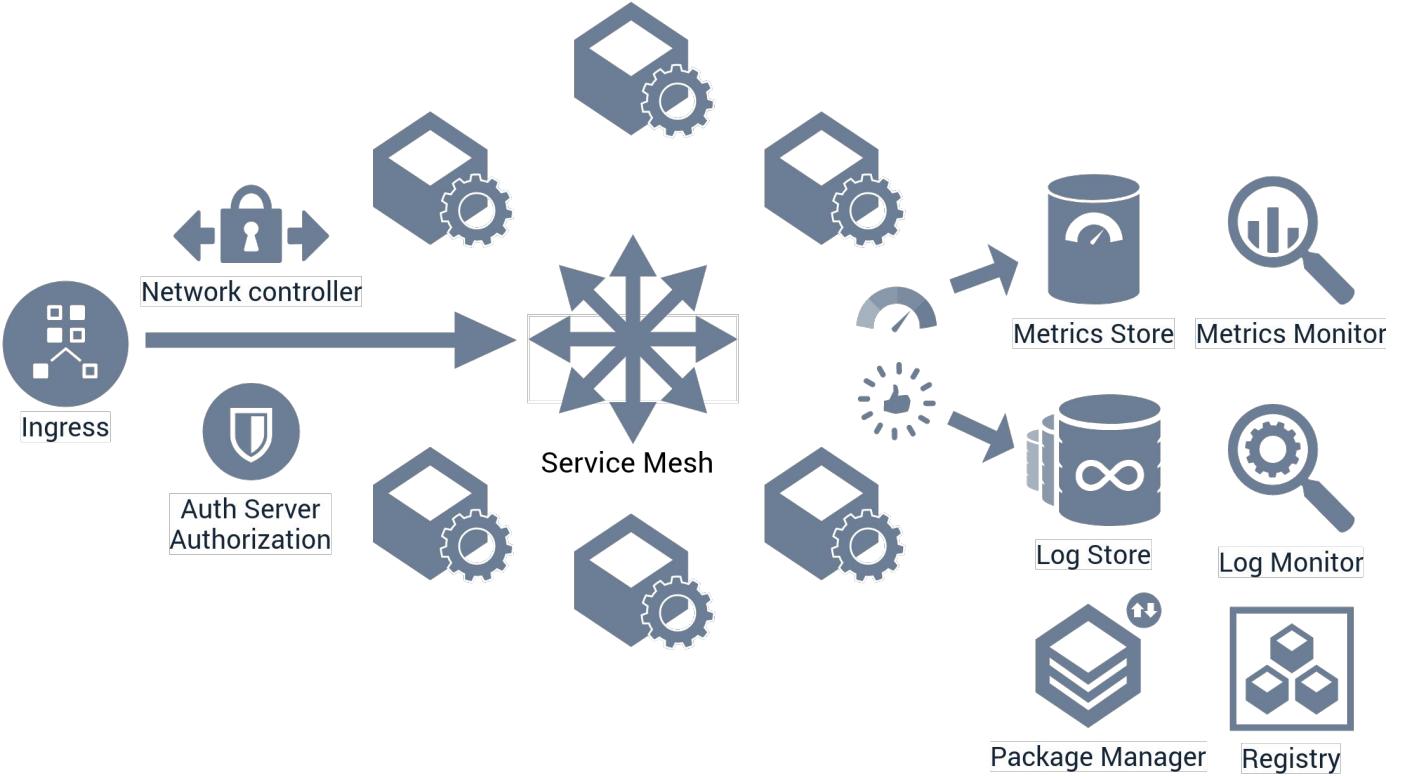














TLS Manager

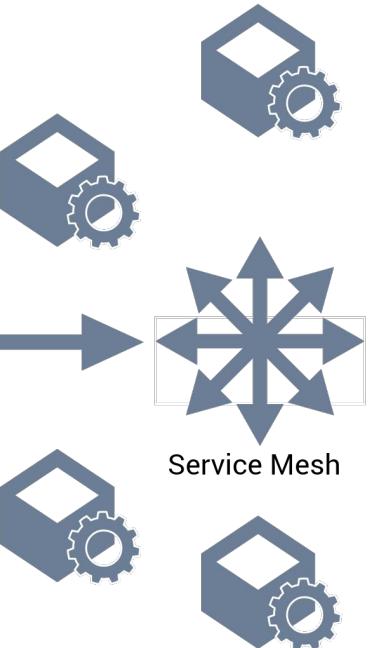


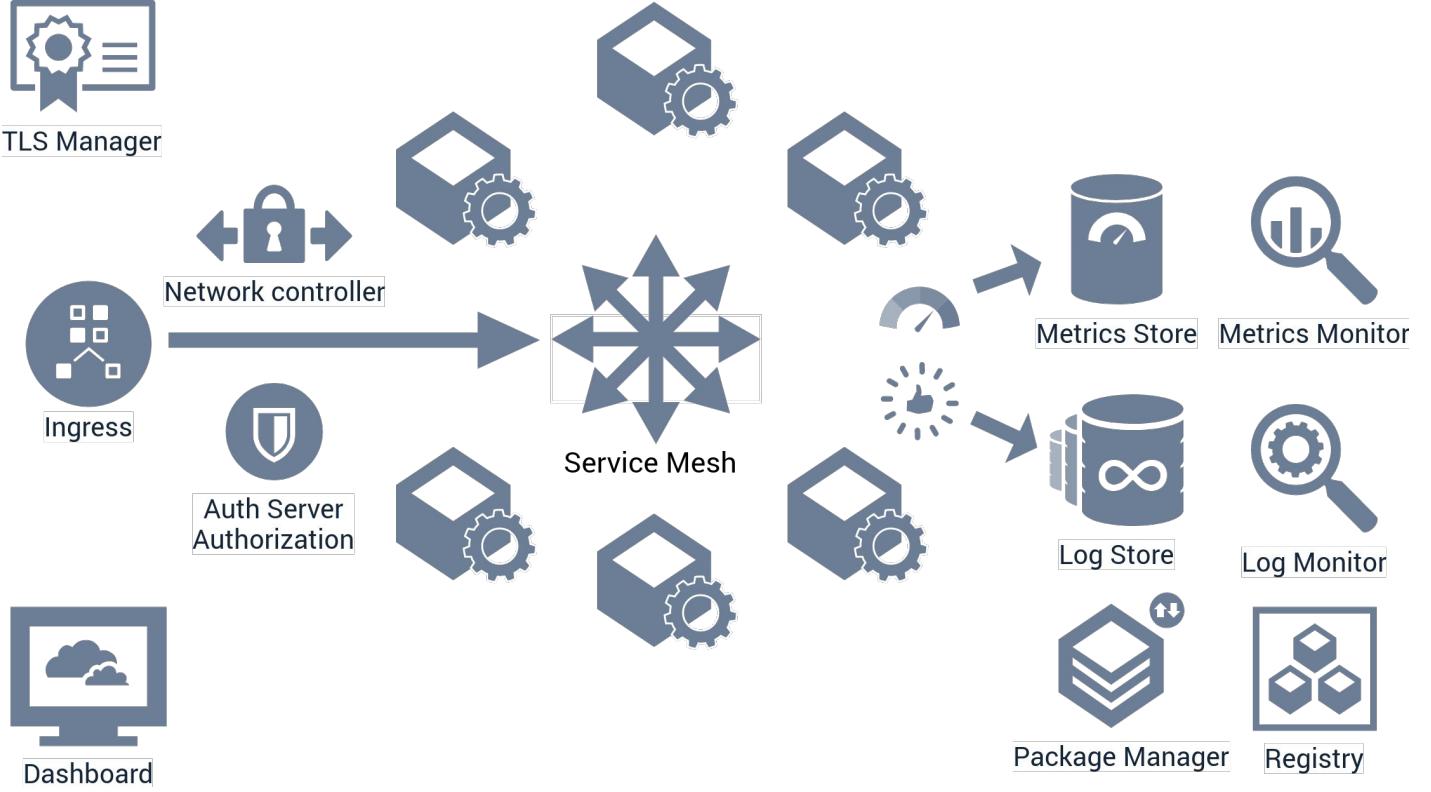
Ingress

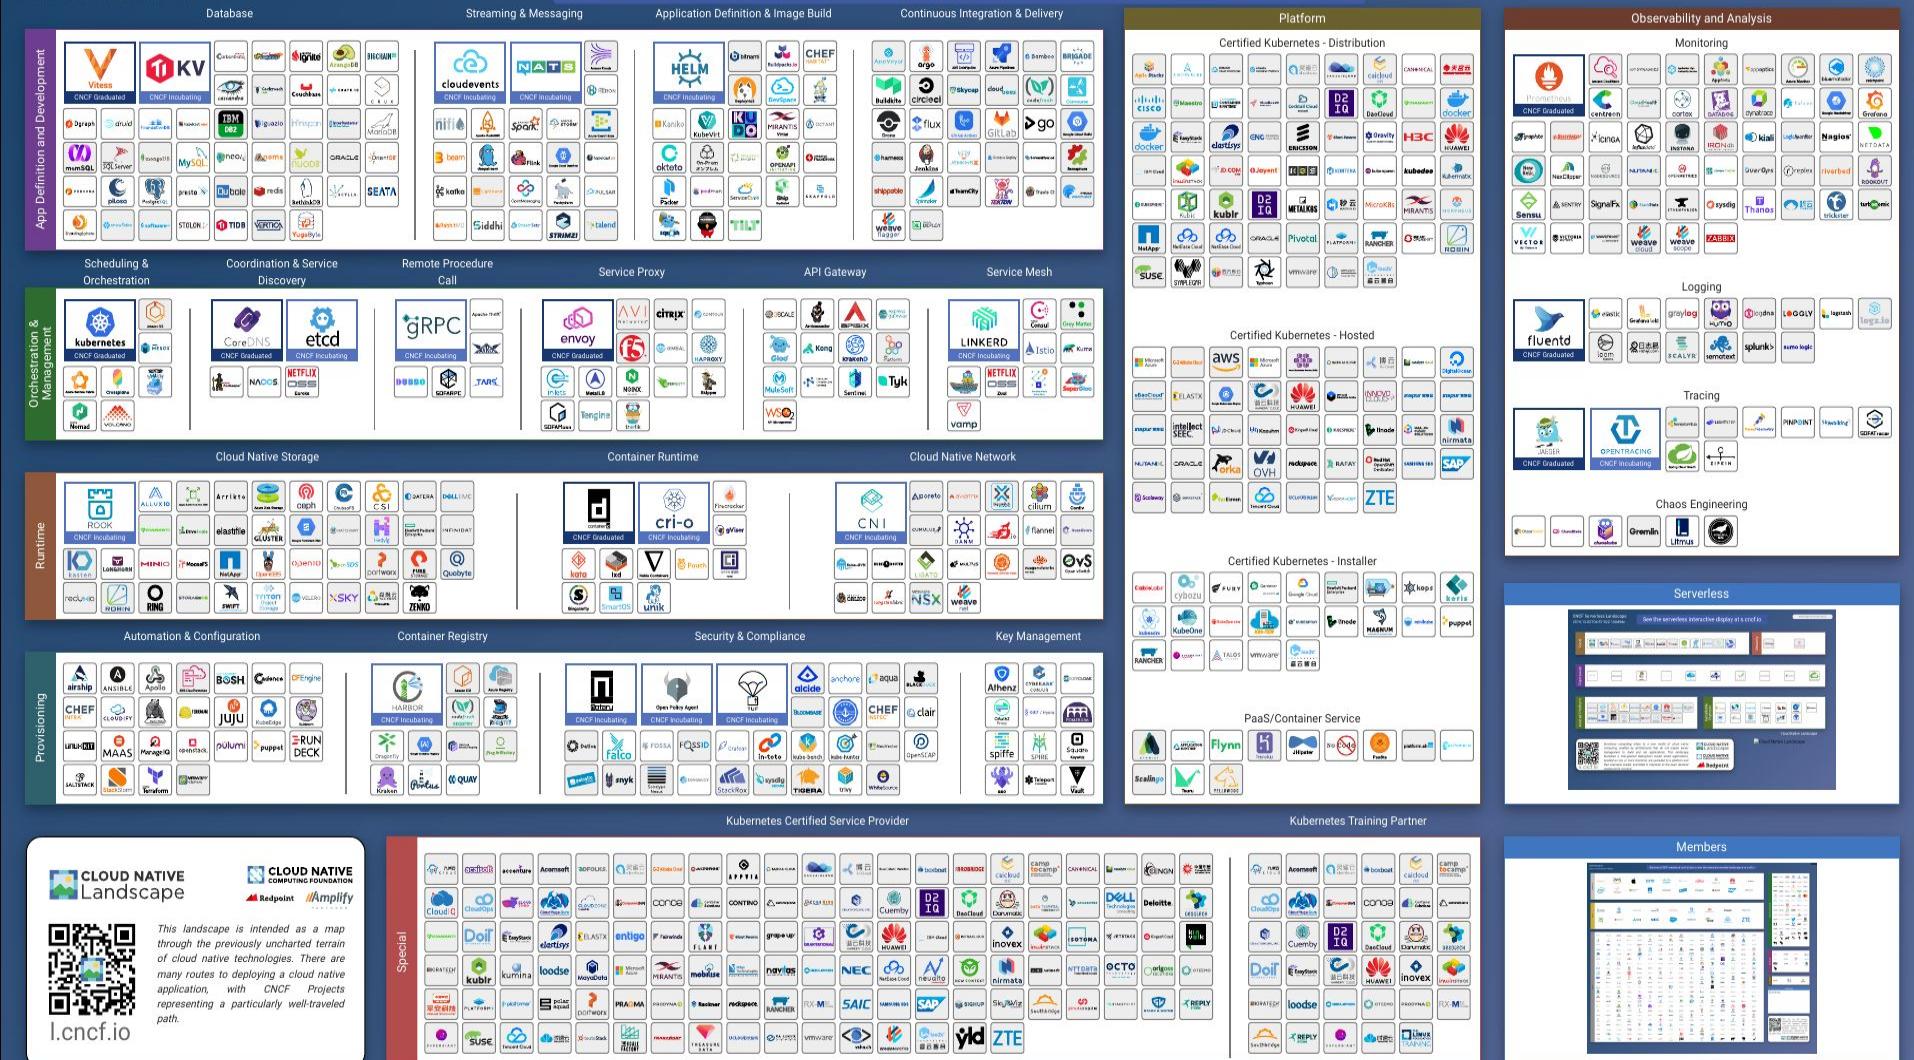
Network controller

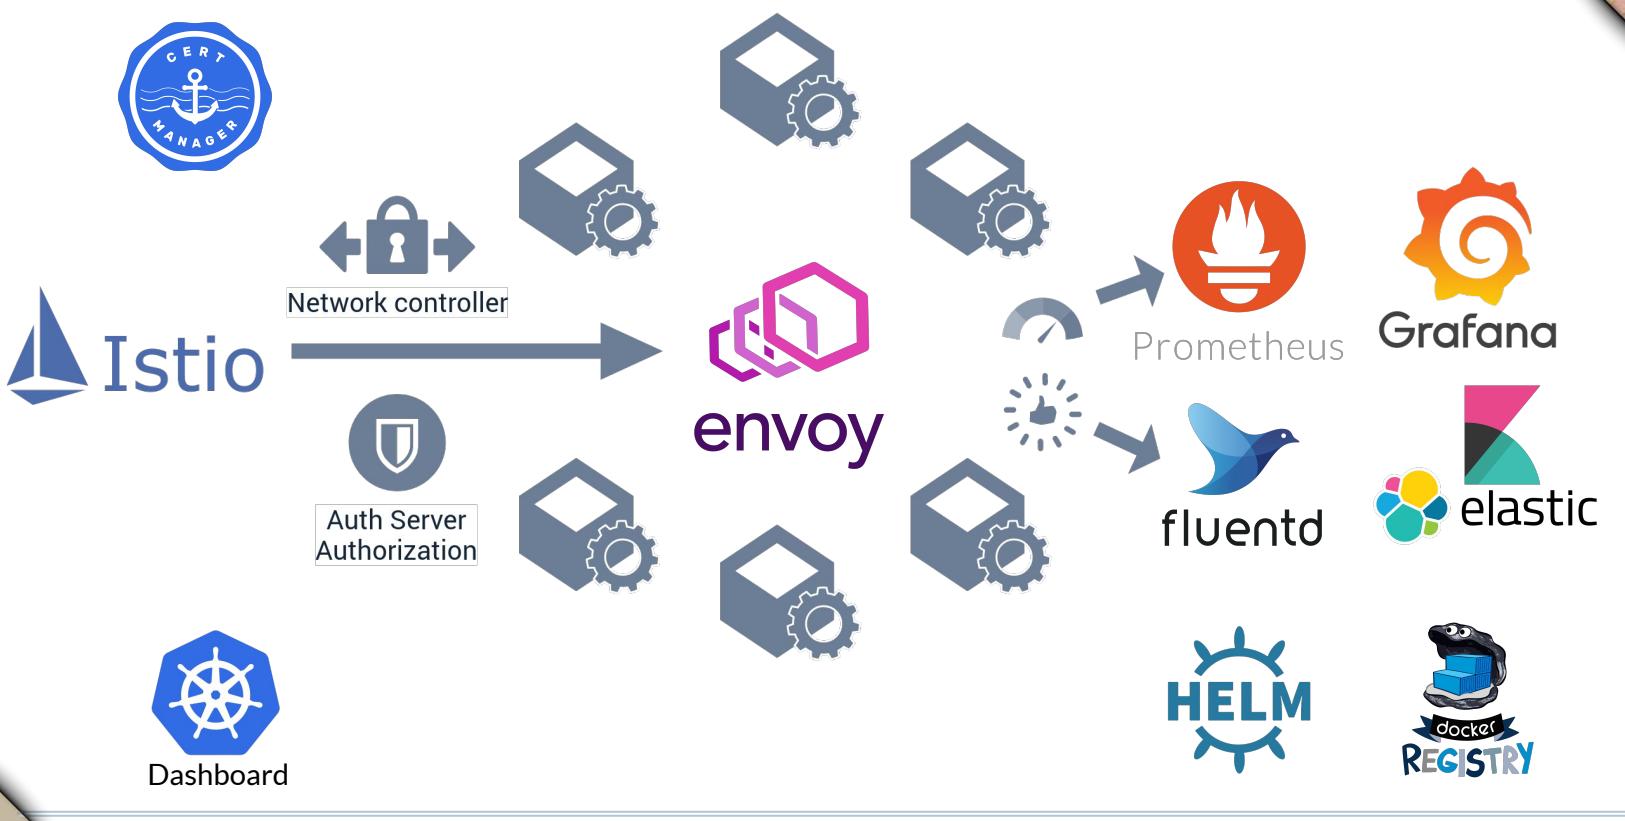


Auth Server
Authorization











cc by 4

@alexivkinx

Declarative deployment

- You can find everything you need in K8s resources
- Kubernetes will do what you need for you
- It does not know what it does not know



No hidden APIs

API's, API's everywhere

- K8s master API, kubelet API, Application APIs
- RBAC and network policies are the only things between you and the cluster controls.



Meet the developer where they are at

- Secrets, config data and the application state are in env variables and folder mounts
- Legacy apps hastily ported to containers
- Security last



Pillage!



cc by 4

@alexivkinx

Pwnage now

Open ports / unauthed access

Supply chain for build layers, OSS Libs, registries

Container escape - Kernel boundary

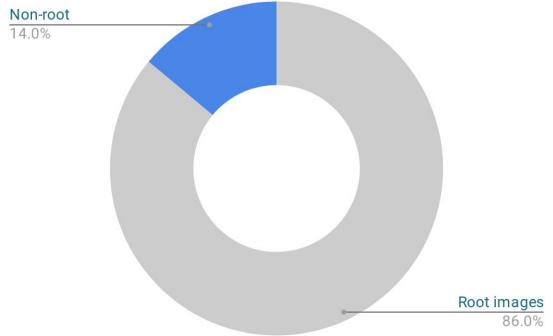
Network sniffing and MITM with eBPF

Node/pod keys

Privesc via secrets/daemonsets

Get a code exec in a pod

- Pentest
 - The app - SSRF is the best
 - The APIs
 -
- Look for known vulns in the k8s apps world



istio / istio

Code Issues 938 Pull requests 84

🔍 is:open label:area/security

✖ Clear current search query, filters, and sorts

167 Open ✓ 791 Closed

grafana / grafana

Code Issues 2,247 Pull requests 125

Pinned issues

Observability Roadmap
#15999 opened on Mar 14 by davikal
Open

Gene #16026 Open

🔍 is:open label:area/backend/security

✖ Clear current search query, filters, and sorts

11 Open ✓ 40 Closed

prometheus / prometheus

Code Issues 364 Pull requests 126

🔍 is:issue is:open

364 Open ✓ 2,656 Closed

envoyproxy / envoy

Code Issues 566 Pull requests 59

🔍 is:open label:area/security

✖ Clear current search query, filters, and sorts

11 Open ✓ 21 Closed

Kubernetes Network Plugins



	Calico	Canal	Cilium	Flannel	Kube-router	WeaveNet
Encryption	No	No	Yes	No	No	Yes
Network Policies	Ingress Egress	Ingress Egress	Ingress Egress	None	Ingress	Ingress Egress

Ingress controllers



	Ingress-nginx	Nginx-ingress	Istio	Ambassador	Traefik	Gloo
Authentication	Yes	No	mTLS, OpenID	Basic, OAuth	Basic, URL	
JWT Support	No		Yes		No	
WAF	ModSecurity		ModSecurity	No	No	

https://docs.google.com/spreadsheets/d/1DnsHtdHbxjvHmxvlu7VhzWcWgLAn_Mc5L1WlhLDA__k/edit#gid=0

Real world Kubernetes

- Security Misconfiguration is the King
 - ➔ Broken Authentication, Access Control
 - ➔ Insufficient Logging and Monitoring
 - ➔ Service accounts/secrets distributions
 - ➔ mTLS and In-transit encryption
- Example - Pod Security Policies are dead



What do?



Easy

- Use simple images to start from or grab and reconstruct dockerfile and build the image yourself
- No privileged containers or shared volumes with the node
- Monitor for rogue containers
- In the cloud - secure metadata
- On baremetal - use hardened OS (Flatcar, Minimal Debian, Alpine with CRI)
- Use RBAC (duh!)



<https://github.com/alexivkin/docker-historian>

Normal

- Build from scratch or use distroless images
- Use PodSecurity Admission Controller
- Have multiple registries and authenticate/authorize access
- Remap root to non-root users for all containers
- Upgrade master nodes and all nodes to the latest version
- SAST/DAST/SCA for images
- Ship out K8s secrets to a vault



Hardmode

- Use Pod Admission Policies
- Aim for zero-trust (mTLS, SPIFFE/SPIRE)
- Sign your images
 - ➔ Use TUF/Notary to address supply chain risks
- Confidential nodes/VM sandboxing
 - ➔ Container separation is weaker than hypervisors
- Unmix sensitive workloads
 - ➔ Namespaces will not provide hard multitenancy
 - ➔ Use separate clusters



Security best practices, frameworks and standards

- CTFs and playgrounds
- CIS Benchmarks for Docker and Kubernetes
- Aqua, StackRox, SysDig, Tigera, Neuvector whitepapers
- Night time reading
 - US DoD DevSecOps reference design
 - NIST SP 800-190 - Compliance in Container Environments
 - SANS Checklist for Audit of Docker containers
- <https://github.com/alexivkin/containerpwn>
- <https://github.com/alexivkin/kubepwn>



**It's gonna get more complex
Errors, not vulnerabilities, lead to a lot of breaches
Make it secure by default**



<https://github.com/alexivkin> 

@alexivkinx 

in/alexivkin 

